

STA414 Assignment 1

Zikun Chen

January 29, 2019

1. Probability and Calculus

1.1 Variance and Covariance

(a) Let $X = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix}$ and $Y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}$

Since they are independent, $\mathbb{E}(x_i y_j) = \mathbb{E}(x_i) \mathbb{E}(y_j)$

By definition,

$$\begin{aligned} Cov(X, Y) &= \mathbb{E}(XY^T) - \mathbb{E}(X) \mathbb{E}(Y)^T \\ &= \mathbb{E}\left(\sum_{i=1}^m x_i y_i\right) - \sum_{i=1}^m \mathbb{E}(x_i) \mathbb{E}(x_j) \\ &= \sum_{i=1}^m \mathbb{E}(x_i y_i) - \sum_{i=1}^m \mathbb{E}(x_i) \mathbb{E}(x_j) \\ &= \sum_{i=1}^m \mathbb{E}(x_i) \mathbb{E}(y_i) - \sum_{i=1}^m \mathbb{E}(x_i) \mathbb{E}(x_j) \\ &= 0 \end{aligned}$$

(b) Let $A = \begin{bmatrix} a_{11} & \dots & a_{1m} \\ \vdots & \ddots & \vdots \\ a_{m1} & \dots & a_{mm} \end{bmatrix}$

$$\begin{aligned}
\mathbb{E}(X + AY) &= \mathbb{E}\left(\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix} + \begin{bmatrix} a_{11} & \dots & a_{1m} \\ \vdots & \ddots & \vdots \\ a_{m1} & \dots & a_{mm} \end{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}\right) \\
&= \mathbb{E}\left(\begin{bmatrix} x_1 + \sum_{j=1}^m y_j a_{1j} \\ \vdots \\ x_i + \sum_{j=1}^m y_j a_{ij} \\ \vdots \\ x_m + \sum_{j=1}^m y_j a_{mj} \end{bmatrix}\right) = \mathbb{E}(X) + \mathbb{E}\left(\begin{bmatrix} \sum_{j=1}^m a_{1j} \mathbb{E}(y_j) \\ \vdots \\ \sum_{j=1}^m a_{ij} \mathbb{E}(y_j) \\ \vdots \\ \sum_{j=1}^m a_{mj} \mathbb{E}(y_j) \end{bmatrix}\right) \\
&= \mathbb{E}(X) + A \mathbb{E}(Y)
\end{aligned}$$

Since X and Y independent and A is constant

$$\begin{aligned}
Var(X + AY) &= Var(X) + Var(AY) \\
&= Var(X) + Cov(AY, AY) \\
&= Var(X) + \mathbb{E}(AY[AY]^T) - E(AY)E(AY)^T \\
&= Var(X) + \mathbb{E}(AYY^T A^T) - AE(Y)(AE(Y))^T \\
&= Var(X) + A \mathbb{E}(YY^T) A^T - AE(Y)E(Y)^T A^T \\
&= Var(X) + A(\mathbb{E}(YY^T) - E(Y)E(Y)^T) A^T \\
&= Var(X) + ACov(Y, Y) A^T \\
&= Var(X) + AVar(Y) A^T
\end{aligned}$$

- (c) Since AX is a linear transformation of X . AX is a Gaussian distribution with mean and covariance matrix as follows:

Since $\mathbb{E}(X) = \mu$ and $Var(X) = \Sigma$

By part b), $\mathbb{E}(AX) = A\mathbb{E}(X) = A\mu$ and $Var(AX) = A\Sigma A^T$

Therefore, $AX \sim \mathcal{N}(A\mu, A\Sigma A^T)$

1.2 Densities

- (a) It is certainly possible that a probability density function takes values greater than 1.
- (b) Normal distribution density function with mean μ and variance σ^2 :

$$\frac{1}{\sigma\sqrt{2\pi}} \exp\left\{-\frac{x - \mu^2}{2\sigma^2}\right\}$$

X be a uni-variate normally distributed random variable with mean 0 and variance 1/100. It's density is of the form:

$$p_X(x) = \frac{1}{\frac{1}{10}\sqrt{2\pi}} \exp\left\{-\frac{x^2}{2\frac{1}{100}}\right\} = \frac{10}{\sqrt{2\pi}} e^{-50x^2} = \sqrt{\frac{50}{\pi}} e^{-50x^2}$$

(c) $p_X(0) = \sqrt{\frac{50}{\pi}}$

(d) $P(X = 0) = \int_0^0 p_X(x) dx = 0$

1.3 Calculus

(a)

$$x^T y = \sum_{i=1}^m x_i y_i$$

$$\frac{\partial x^T y}{\partial x} = \left(\frac{\partial x^T y}{\partial x_1}, \dots, \frac{\partial x^T y}{\partial x_m} \right)^T = (y_1, \dots, y_m)^T = y$$

(b)

$$x^T x = \sum_{i=1}^m x_i^2$$

$$\frac{\partial x^T x}{\partial x} = \left(\frac{\partial x^T x}{\partial x_1}, \dots, \frac{\partial x^T x}{\partial x_m} \right)^T = (2x_1, \dots, 2x_m)^T = 2x$$

(c)

$$x^T A x = \sum_{i=1}^m \sum_{j=1}^m a_{ij} x_i x_j$$

For a particular $k \in \mathbb{Z}, k \in [1, m]$

$$\frac{\partial x^T A x}{\partial x_k} = \sum_{i=1}^m a_{ik} x_i x_k + \sum_{j=1}^m a_{kj} x_k x_j$$

$$\frac{\partial x^T A x}{\partial x} = \left(\frac{\partial x^T A x}{\partial x_1}, \dots, \frac{\partial x^T A x}{\partial x_m} \right)^T = \left(\sum_{i=1}^m (a_{1i} + a_{i1}) x_i, \dots, \sum_{i=1}^m (a_{mi} + a_{im}) x_i \right)^T = (A + A^T) x$$

(d)

$$A x = \begin{bmatrix} \sum_{i=1}^m a_{1i} x_i \\ \vdots \\ \sum_{i=1}^m a_{mi} x_i \end{bmatrix}$$

$$\frac{\partial A x}{\partial x} = \begin{bmatrix} a_{11} & \dots & a_{1m} \\ \vdots & \ddots & \vdots \\ a_{m1} & \dots & a_{mm} \end{bmatrix} = A$$

2. Regression

2.1 Linear Regression

- (a) **Distribution:** Since $(X^T X)^{-1} X^T$ is a matrix, $\hat{\beta} = (X^T X)^{-1} X^T Y$ is a linear transformation of Y. Therefore, $\hat{\beta}$ follows Normal distribution with the following mean and variance.

Mean:

$$\mathbb{E}(\hat{\beta}) = \mathbb{E}((X^T X)^{-1} X^T Y) = (X^T X)^{-1} X^T \mathbb{E}(Y) = (X^T X)^{-1} X^T X \beta = \beta$$

Variance:

$$\text{Let } A = (X^T X)^{-1} X^T$$

$$\begin{aligned} \text{Var}(\hat{\beta}) &= \text{Var}((X^T X)^{-1} X^T Y) = \text{Var}(AY) = A \text{Var}(Y) A^T \\ &= (X^T X)^{-1} X^T \sigma^2 I ((X^T X)^{-1} X^T)^T \\ &= \sigma^2 (X^T X)^{-1} X^T (X ((X^T X)^{-1})^T) \\ &= \sigma^2 (X^T X)^{-1} X^T X ((X^T X)^T)^{-1} \\ &= \sigma^2 (X^T X)^{-1} X^T X ((X^T X))^{-1} \\ &= \sigma^2 (X^T X)^{-1} \end{aligned}$$

$$\text{Therefore, } \hat{\beta} \sim \mathcal{N}(\beta, \sigma^2 (X^T X)^{-1})$$

- (b) Likelihood Function:

$$\mathcal{L}(\beta|Y) = -\frac{m}{2} \ln(2\pi) - \frac{1}{2} \ln(|\Sigma|) - \frac{\sigma^2}{2} (Y - X\beta)^T (Y - X\beta)$$

$$\text{Set } \frac{\partial \log(\mathcal{L}(\beta|Y))}{\partial \beta} = 0$$

$$\frac{\partial \log(\mathcal{L}(\beta|Y))}{\partial \beta} = \frac{\partial [-\frac{\sigma^2}{2} (Y^T Y - 2\beta^T X^T Y + \beta^T X^T X \beta)]}{\partial \beta} = -\frac{\sigma^2}{2} [-2X^T Y + 2(X^T X)\beta] = 0$$

$$-X^T Y + (X^T X)\beta = 0$$

$$\hat{\beta}_{ML} = (X^T X)^{-1} X^T Y$$

- (c) For each i from 1 to m, the marginal distribution of $\hat{\beta}_i$ is:

$$\hat{\beta}_i \sim \mathcal{N}(\beta_i, \sigma^2 C_{ii})$$

where C_{ii} is the ith diagonal element of $(X^T X)^{-1}$

Consequently,

$$\frac{\hat{\beta}_i - \beta_i}{\sigma\sqrt{C_{ii}}} \sim \mathcal{N}(0, 1)$$

$$\begin{aligned} \mathbb{P}(|\hat{\beta}_i - \beta_i| \leq \epsilon) &= \mathbb{P}(\beta_i - \epsilon \leq \hat{\beta}_i \leq \beta_i + \epsilon) \\ &= \mathbb{P}\left(\frac{-\epsilon}{\sigma\sqrt{C_{ii}}} \leq \frac{\hat{\beta}_i - \beta_i}{\sigma\sqrt{C_{ii}}} \leq \frac{\epsilon}{\sigma\sqrt{C_{ii}}}\right) = \Phi\left(\frac{\epsilon}{\sigma\sqrt{C_{ii}}}\right) - \Phi\left(\frac{-\epsilon}{\sigma\sqrt{C_{ii}}}\right) \end{aligned}$$

2.2 Ridge Regression

(a) $\hat{\beta}_{MAP}$ estimator maximize the posterior $p(\beta|X, Y) \propto p(X, Y|\beta)P(\beta)$

$$p(X, Y|\beta) \propto \mathcal{L}(\beta|Y) \propto \exp\left\{-\frac{\sigma^2}{2}(Y - X\beta)^T(Y - X\beta)\right\}$$

$$P(\beta) \propto \exp\left\{\frac{\tau^2}{2}\beta^T\beta\right\}$$

Set

$$\frac{\partial \log(p(X, Y|\beta)p(\beta))}{\partial \beta} = 0$$

$$\log[p(X, Y|\beta)] + \log[p(\beta)] \propto -\frac{\sigma^2}{2}(Y^T Y - 2\beta^T X^T Y + \beta^T X^T X \beta) - \frac{\tau^2}{2}\beta^T \beta$$

$$\frac{\partial \log(p(X, Y|\beta)p(\beta))}{\partial \beta} = \frac{\sigma^2}{2}[-2X^T Y + 2(X^T X)\beta] + \tau^2 \beta = -\sigma^2 X^T Y + \sigma(X^T X)\beta + \tau^2 \beta = 0$$

$$\sigma^2(X^T X)\beta + \tau^2 \beta = \sigma^2 X^T Y$$

$$(X^T X)\beta + \frac{\tau^2}{\sigma^2}\beta = X^T Y$$

$$(X^T X + \frac{\tau^2}{\sigma^2}I)\beta = X^T Y$$

$$\hat{\beta}_{MAP} = (X^T X + \lambda I)^{-1} X^T Y$$

$$(b) \text{ Let } X^* = \begin{bmatrix} x_{11} & \dots & x_{1m} \\ \vdots & \ddots & \vdots \\ x_{m1} & \dots & x_{mm} \\ \sqrt{\lambda} & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \sqrt{\lambda} \end{bmatrix}$$

$$\text{Then } X^{*T}X^* = \begin{bmatrix} x_{11} & \dots & x_{n1} & \sqrt{\lambda} & \dots & 0 \\ \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\ x_{1m} & \dots & x_{nm} & 0 & \dots & \sqrt{\lambda} \end{bmatrix} \begin{bmatrix} x_{11} & \dots & x_{1m} \\ \vdots & \ddots & \vdots \\ x_{n1} & \dots & x_{nm} \\ \sqrt{\lambda} & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \sqrt{\lambda} \end{bmatrix}$$

$$= \begin{bmatrix} \sum_{i=1}^m x_{i1}^2 + \lambda & \dots & \sum_{j=1}^m x_{i1}x_{jm} \\ \vdots & \ddots & \vdots \\ \sum_{i=1}^m x_{im}x_{i1} & \dots & \sum_{j=1}^m x_{im}x_{jm} + \lambda \end{bmatrix} = X^T X + \lambda I$$

$$X^{*T}Y = \begin{bmatrix} x_{11} & \dots & x_{n1} & \sqrt{\lambda} & \dots & 0 \\ \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\ x_{1m} & \dots & x_{nm} & 0 & \dots & \sqrt{\lambda} \end{bmatrix} \begin{bmatrix} y_1 \\ \vdots \\ y_m \\ 0 \\ \vdots \\ 0 \end{bmatrix} = X^T Y$$

Therefore, MLE estimate of β based on X^* and Y^* is

$$\beta_{ML}^* = (X^{*T}X^*)^{-1}X^{*T}Y = (X^T X + \lambda I)^{-1}X^T Y$$

2.3 Cross Validation

(a) See code in Appendix

```
### Part (a) ###
# Load data from dataset.mat
dataset = ...
```

(b) See 6 functions in Appendix

```
### Part (b) ###
def shuffle_data(data):
    ...
```

(c) See code in Appendix

```


### Part (c) ###
def lambda_train_test(train_data, test_data, lambda_seq)
    ...

```

- (d) In the program, for reproducibility purposes, I set the random seed to 2019. As we observe in the following graph, rolling through different λ in the cross validation process. The best choice of λ is obtained at $\lambda^* = 0.542353$, which minimizes the test error, 5-fold validation error and 10-fold validation error.

The training error increases as we increase the value of λ . The test error curve, 5-fold validation error curve, and 10-fold validation error curve have all have steep drops as we initially increase the value of lambda starting from 0.02. The curves then decrease in a slower pace, eventually hitting the best λ . The curves slope upwards after the best λ .

Additionally, 10-fold validation error achieved better error rate than 5-fold validation.



Test error is minimized when $\lambda = 0.542353$
 5-fold validation error is minimized when $\lambda = 0.542353$
 10-fold validation error is minimized when $\lambda = 0.542353$

Appendix: Python Code

```
"""
STA414 Assignment 1
Due Jan 29, 2019

@name: cross_validation.py
@author: zikunchen
"""

import scipy.io as sio
import numpy as np
import matplotlib.pyplot as plt

### Part (b) ###

def shuffle_data(data):
    """Returns randomly uniformly permuted version of data along the samples.
    Note that y and X need to be permuted the
    same way preserving the target-feature pairs."""
    indices = np.arange(len(data['y']))
    np.random.seed(SEED)

    indices = np.random.permutation(indices)

    return {'y': np.array(data['y'])[indices].tolist(), \
            'X': np.array(data['X'])[indices].tolist()}

def split_data(data, num_folds, fold):
    """
    - Attributes -
    data:
        (y, X) pair in the form of dictionary
    num_folds:
        number of folds in total
    fold:
        specific fold number to be used as the training set

    - Returns -
    data_fold:
        (y, X) pair used for validation
    data_rest:
        rest of the data used for training
    """

    N = len(data['y'])
```



```

fold_size = int(N//num_folds)

fold_indices = np.arange((fold - 1) * fold_size, fold * fold_size)
data_fold = {'y': np.array(data['y'])[fold_indices].tolist(), \
              'X': np.array(data['X'])[fold_indices].tolist()}

rest_indices = np.append(np.arange(0, (fold - 1) * fold_size), \
                          np.arange(fold * fold_size, N), axis = 0)
data_rest = {'y': np.array(data['y'])[rest_indices].tolist(), \
             'X': np.array(data['X'])[rest_indices].tolist()}

return data_fold, data_rest

def train_model(data, lambd):
    """ Returns the coefficients of ridge regression with penalty level ."""

    N = len(data['y'])

    X = np.array(data['X'])
    y = np.array(data['y']).reshape(N,1)

    M = X.shape[1]

    XTX = np.dot(X.T, X)

    model = np.dot(np.dot(np.linalg.inv(XTX + lambd*np.identity(M)), X.T), y)

    return np.squeeze(model).tolist()

def predict(data, model):
    """ Returns the predictions based on data and model."""

    X = np.array(data['X'])
    M = X.shape[1]
    beta = np.array(model).reshape(M,1)

    return np.squeeze(np.dot(X, beta)).tolist()

def loss(data, model):
    """Returns the average squared error loss based on model."""

    N = len(data['y'])
    y = np.array(data['y']).reshape(N,1)

```

```

    predictions = np.array(predict(data, model)).reshape(N, 1)

    error = np.sum(np.squeeze((y - predictions)**2)) / N
    return error

def cross_validation(data, num_folds, lambd_seq):
    """Returns the cross validation error across all lambdas in lambd_seq"""
    data = shuffle_data(data)
    cv_error = [0] * len(lambd_seq)

    for i in range(len(lambd_seq)):
        lambd = lambd_seq[i]
        cv_loss_lmd = 0

        for fold in range(1, num_folds + 1):
            val_cv, train_cv = split_data(data, num_folds, fold)
            model = train_model(train_cv, lambd)
            cv_loss_lmd += loss(val_cv, model)
        cv_error[i] = cv_loss_lmd/num_folds

    return cv_error

### Part (c) ###

def lambd_train_test(train_data, test_data, lambd_seq):
    """Returns the training and test error across all lambdas in lambd_seq"""

    train_error = [0] * len(lambd_seq)
    test_error = [0] * len(lambd_seq)

    for i in range(len(lambd_seq)):
        lambd = lambd_seq[i]
        model = train_model(train_data, lambd)
        train_error[i] = loss(train_data, model)
        test_error[i] = loss(test_data, model)

    return train_error, test_error

### Helper Function ###

def find_min_lambd(error_list):
    """Find the best lambda based on the minimum error achieved in error_list"""
    min_error, index = min((val, idx) for (idx, val) in enumerate(five_fold_cv))
    return lambd_seq[index]

```

```

if __name__ == '__main__':

    ### Part (a) ###
    dataset = sio.loadmat("/Users/zikunchen/Desktop/STA414/A1/dataset.mat")
    data_train_X = dataset["data_train_X"]
    data_train_y = dataset["data_train_y"][0]
    data_test_X = dataset["data_test_X"]
    data_test_y = dataset["data_test_y"][0]

    # Set Random Seed to reproduce results
    SEED = 2019

    # Construct datasets
    train_data = {'y': data_train_y, 'X': data_train_X}
    test_data = {'y': data_test_y, 'X': data_test_X}

    # Construct Lambda values
    space = (1.5-0.02)/(50-1+2)
    lambd_seq = np.arange(0.02, 1.5, space).tolist()

    # Compute training, validation and test errors
    train_error, test_error = lambd_train_test(train_data, test_data, lambd_seq)
    five_fold_cv = cross_validation(train_data, 5, lambd_seq)
    ten_fold_cv = cross_validation(train_data, 10, lambd_seq)

    #Find the lambda with the least error
    lambd_test = find_min_lambd(test_error)
    lambd_five_fold= find_min_lambd(five_fold_cv)
    lambd_ten_fold = find_min_lambd(ten_fold_cv)

    # Plot 'Training Error', 'Test Error',
    # '5-fold Validation Error', and '10-fold Validation Error'
    # curves against Lambda values
    fig, ax = plt.subplots()

    test_min = min(test_error)
    xpos = test_error.index(test_min)
    xmin = lambd_seq[xpos]

    ax.annotate('Best lambda = %f' % lambd_test, xy=(xmin, test_min), \
                xytext=(xmin, test_min+1), \
                arrowprops=dict(facecolor='red', shrink=0.05),)

    plt.plot(lambd_seq, train_error, 'mo-')
    plt.plot(lambd_seq, test_error, 'bo-')
    plt.plot(lambd_seq, five_fold_cv, 'go-')

```

```

plt.plot(lambd_seq, ten_fold_cv, 'yo-')

plt.title('Errors vs. Lambda')
plt.ylabel('Error')
plt.xlabel('')
plt.legend(['Training Error', 'Test Error', '5-fold Validation Error', \
           '10-fold Validation Error'], loc='upper right')

plt.show()

print("Test error is minimized when lambda = %f" % lambd_test)
print("5-fold validation error is minimized when lambda = %f" % lambd_five_fold)
print("10-fold validation error is minimized when lambda = %f" % lambd_ten_fold)

```