

STA414 - Probabilistic Learning


Assignment 3
Sampling and Gradient Estimators

Zikun Chen
1001117882
April 5, 2019

1 Nature's Rejection Sampler

(a) Numerical Integration

Unnormalized probability $p(x, g = 0 | \theta = 0)$, shown below as a function of x from -20 to 20:



Using the quadrature function, we estimate that $p(g = 0 | \theta = 0) = 0.803$


(b) Rejection Sampling

Choose proposal normal density $q(x) = p(x | \theta = 0)$ and the unnormalized probability $p(x) = p(g = 1, x | \theta = 0)$.

Furthermore, it is also appropriate to use $M = 1$ to be the bound for $\frac{p(x)}{q(x)}$, since:

$$\frac{p(x)}{q(x)} = \frac{p(g = 1, x | \theta = 0)}{p(x | \theta = 0)} = p(g = 1 | x, \theta = 0) \leq 1$$

Histogram of 10000 accepted samples:



Rejection Sampling:
 Fraction of accepted sample is 0.803

(c) Importance Sampling

We want to estimate:

$$p(g = 0 | \theta = 0) = \int p(g = 0, x | \theta = 0) p(x | \theta = 0) dx$$


In this case, $f(x) = p(g = 0, x | \theta = 0)$, our target distribution is $\tilde{p}(x) = p(x | \theta = 0)$
 This is the same as the proposal density $q(x) = p(x | \theta = 0)$
 Therefore:

$$\begin{aligned} \hat{e}(x_1, x_2, \dots, x_K) &= \frac{1}{K} \sum_{i=1}^K f(x_i) \frac{\frac{\tilde{p}(x_i)}{q(x_i)}}{\frac{1}{K} \sum_{j=1}^K \frac{\tilde{p}(x_j)}{q(x_j)}} = \sum_{i=1}^K f(x_i) \frac{1}{\sum_{j=1}^K 1} = \frac{1}{K} \sum_{i=1}^K f(x_i) \\ &= \frac{1}{K} \sum_{i=1}^K p(g = 0, x_i | \theta = 0) = \frac{1}{K} \sum_{i=1}^K \left[1 - \frac{\sin^2(5(x_i))}{25 \sin^2(x_i)} \right] \quad \text{where } x_i \sim_{\text{iid}} N(0, 4^2) \end{aligned}$$

This is the same as a simple monte carlo sampling of $p(g = 0, x_i | \theta = 0)$

Importance Sampling:
 Estimate of fraction of photons that get absorbed is 0.802

(d) Plot unnormalized density $p(x = 1.7, g = 1, \theta)$




(e) Metropolis Hasting

Implementation:

```
def metro_hast(prop_sigma, N, seed, sigma=4, x=1.7, g=1):
    np.random.seed(seed)
    t = np.random.uniform(-20, 20, 1)[0]
    i = 0
    count = 0
    thetas = np.empty(N)
    while i < N:
        tnew = np.random.normal(t, prop_sigma, 1)[0]
        gtnew_t = st.norm.pdf(tnew, loc=t, scale=prop_sigma)
        gt_tnew = st.norm.pdf(t, loc=tnew, scale=prop_sigma)
        pt = p_theta(t, x, g)
        ptnew = p_theta(tnew, x, g)
        accept = min(1, (ptnew * gt_tnew) / (pt * gtnew_t))
        u = np.random.uniform(0, 1, 1)[0]
        if u < accept:
            t = tnew
            thetas[i] = t
            i += 1
            count += 1
    return thetas, acceptance
```

After experimenting with a combination of different standard deviations and sample sizes, a standard deviation of 3 is chosen for the proposal distribution and 50000 samples are drawn, which are shown below:



(f) Estimate posterior probability from samples

$$p(-3 < \theta < 3 | x = 1.7, g = 1) = \int_{-3}^3 p(\theta | x = 1.7, g = 1) d\theta$$

We calculate the proportion of values between -3 and 3 in the samples that we drew from part e):

```
p_abs_three = len(sample_t[(sample_t<3) & (sample_t>-3)]) \
               / len(sample_t)

>>> Posterior probability of theta between +3 and -3 is 0.520
```

2 Gradient Estimators

(a) Score Function has Zero Expectation

$$\begin{aligned} \mathbb{E}_{p(x|\theta)} [\nabla_{\theta} \log p(x|\theta)] &= \mathbb{E}_{p(x|\theta)} \left[\frac{\nabla_{\theta} p(x|\theta)}{p(x|\theta)} \right] \\ &= \int p(x|\theta) \frac{\nabla_{\theta} p(x|\theta)}{p(x|\theta)} dx \\ &= \nabla_{\theta} \int p(x|\theta) dx = \nabla_{\theta} 1 = 0 \end{aligned}$$

(b) REINFORCE is unbiased

$$\begin{aligned}
\mathbb{E}_{p(b|\theta)} \left[f(b) \frac{\partial}{\partial \theta} \log p(b|\theta) \right] &= \mathbb{E}_{p(b|\theta)} \left[f(b) \frac{\frac{\partial}{\partial \theta} p(b|\theta)}{p(b|\theta)} \right] \\
&= \int f(b) \frac{\partial}{\partial \theta} p(b|\theta) db \\
&= \frac{\partial}{\partial \theta} \int f(b) p(b|\theta) db = \frac{\partial}{\partial \theta} \mathbb{E}_{p(b|\theta)} [f(b)]
\end{aligned}$$

(c) REINFORCE with fixed baseline is unbiased

Show the estimator is unbiased:

$$\begin{aligned}
\mathbb{E}_{p(b|\theta)} \left[c \frac{\partial}{\partial \theta} \log p(b|\theta) \right] &= c \mathbb{E}_{p(b|\theta)} \left[\frac{\frac{\partial}{\partial \theta} p(b|\theta)}{p(b|\theta)} \right] \\
&= c \int \frac{\partial}{\partial \theta} p(b|\theta) db \\
&= c \frac{\partial}{\partial \theta} \int p(b|\theta) db = 0
\end{aligned}$$

Therefore,

$$\mathbb{E}_{p(b|\theta)} \left[[f(b) - c] \frac{\partial}{\partial \theta} \log p(b|\theta) \right] = \mathbb{E}_{p(b|\theta)} \left[f(b) \frac{\partial}{\partial \theta} \log p(b|\theta) \right] - \mathbb{E}_{p(b|\theta)} \left[c \frac{\partial}{\partial \theta} \log p(b|\theta) \right] = 0$$

(d) Biased REINFORCE with dependant baseline

Show estimator is unbiased:

$$\begin{aligned}
\mathbb{E}_{p(b|\theta)} \left[[f(b) - c(b)] \frac{\partial}{\partial \theta} \log p(b|\theta) \right] &= \mathbb{E}_{p(b|\theta)} \left[f(b) \frac{\partial}{\partial \theta} \log p(b|\theta) \right] - \mathbb{E}_{p(b|\theta)} \left[c(b) \frac{\partial}{\partial \theta} \log p(b|\theta) \right] \\
&\text{by part b)} = \frac{\partial}{\partial \theta} \mathbb{E}_{p(b|\theta)} [f(b)] - \frac{\partial}{\partial \theta} \mathbb{E}_{p(b|\theta)} [c(b)] \\
&\neq \frac{\partial}{\partial \theta} \mathbb{E}_{p(b|\theta)} [f(b)]
\end{aligned}$$

3 Comparing variances of gradient estimators

(a) Variance of Monte Carlo Estimator

$$\mathbb{V} [\hat{L}_{MC}] = \mathbb{V} \left[\sum_{i=1}^D x_d \right] = \sum_{i=1}^D \mathbb{V} [x_d] + \sum_{i \neq j} Cov(x_i, x_j) = \sum_{i=1}^D 1 = D$$

(b) REINFORCE estimator with a baseline

Let $x = [x_1, \dots, x_n]^T$ and $\theta = [\theta_1, \dots, \theta_n]^T$

$$\log p(x|\theta) = \log \frac{1}{\sqrt{2\pi^D|I|}} \exp\left\{-\frac{1}{2}(x-\theta)^T(x-\theta)\right\}$$

$$\begin{aligned} \hat{g}^{\text{SF}}[f] &= [f(x) - c(\theta)] \frac{\partial}{\partial \theta} \log p(x|\theta) \\ &= \left[\sum_{d=1}^D x_d - \sum_{d=1}^D \theta_d \right] \nabla_{\theta} \log \frac{1}{\sqrt{2\pi^D|I|}} \exp\left\{-\frac{1}{2}(x-\theta)^T(x-\theta)\right\} \\ &= \left[\sum_{d=1}^D (x_d - \theta_d) \right] \left(-\frac{1}{2}\right) \nabla_{\theta} (x-\theta)^T(x-\theta) \\ &= \left[\sum_{d=1}^D (x_d - \theta_d) \right] \left(-\frac{1}{2}\right)(-2)(x-\theta) \\ &= \left[\sum_{d=1}^D (x_d - \theta_d) \right] (x-\theta) \\ &= \epsilon(\epsilon^T \mathbb{1}) \end{aligned}$$

where $\epsilon = [\epsilon_1, \dots, \epsilon_D]^T \sim N(0, I)$ and $\mathbb{1} = [1, \dots, 1]^T$

(c) Variance of REINFORCE estimator

For any $d \in \{1, \dots, D\}$:

$$E[\epsilon_d^2] = \mathbb{V}[\epsilon_d] + E[\epsilon_d]^2 = 1$$

$$\begin{aligned}
\mathbb{V} [\hat{g}_1^{\text{SF}}] &= \mathbb{V} [\epsilon_1(\epsilon^T \mathbb{1})] \\
&= E [\epsilon_1^2(\epsilon^T \mathbb{1})^2] - E [\epsilon_1(\epsilon^T \mathbb{1})]^2 \\
&= E \left[\epsilon_1^2 \left(\sum_{d=1}^D \epsilon_d \right)^2 \right] - E \left[\epsilon_1 \sum_{d=1}^D \epsilon_d \right]^2 \\
&= E \left[\epsilon_1^2 \left(\sum_{d=1}^D \epsilon_d^2 + 2 \sum_{i < j}^D \epsilon_i \epsilon_j \right) \right] - E \left[\epsilon_1^2 + \epsilon_1 \sum_{d=2}^D \epsilon_d \right]^2 \\
&= E [\epsilon_1^4] + E \left[\epsilon_1^2 \sum_{d=2}^D \epsilon_d^2 \right] + 2 \sum_{i < j}^D E [\epsilon_i \epsilon_j] - (E [\epsilon_1^2] + 0)^2 \\
&= 3 + E [\epsilon_1^2] E \left[\sum_{d=2}^D \epsilon_d^2 \right] + 0 - 1^2 \\
&= 3 + (D-1) - 1 = D+1
\end{aligned}$$

(d) Reparameterization of gradient estimator

$$\begin{aligned}
\nabla_{\theta} L(\theta) &= \nabla_{\theta} E_{x \sim p(x|\theta)} [f(x)] \\
&= \nabla_{\theta} E_{x \sim p(\epsilon)} [f(T(\theta, \epsilon))] \\
&= E_{x \sim p(\epsilon)} [\nabla_{\theta} f(T(\theta, \epsilon))] \\
&= E_{x \sim p(\epsilon)} [\nabla_{\theta} f(\theta + \epsilon)] \\
&= E_{x \sim p(\epsilon)} \left[\nabla_{\theta} \sum_{d=1}^D [\theta_d + \epsilon_d] \right] \\
&= E_{x \sim p(\epsilon)} [\nabla_{\theta} \mathbb{1}] \\
&= E_{x \sim p(\epsilon)} [\mathbb{1}] = \mathbb{1}
\end{aligned}$$

$$\hat{g}^{\text{REPARAM}} = \frac{\partial f}{\partial x} \frac{\partial x}{\partial \theta} = \mathbb{1}^T I = \mathbb{1}^T$$

$$\mathbb{V} [\hat{g}_1^{\text{REPARAM}}] = \mathbb{V} [1] = 0 < \mathbb{V} [\hat{g}_1^{\text{SF}}]$$

Again $\mathbb{1} = [1, \dots, 1]^T$

4 Appendix: Code

```

import numpy as np
import seaborn as sns
import scipy.stats as st
import scipy
import matplotlib.pyplot as plt

```



```

sns.set()

def pg(x, g, theta=0, sigma=4):
    result = scipy.sin(5*(x-theta))**2/(25*scipy.sin(x-theta)**2)
    return result if g == 1 else (1-result)

def px(x, theta=0, sigma=4):
    return st.norm.pdf(x, loc=theta, scale=sigma)

def smc(g, N, seed, theta=0, sigma=4):
    np.random.seed(SEED)
    X = np.random.normal(theta, sigma, N)
    return np.sum([pg(x, g) for x in X])/N

def rejection_sampler(M, N, seed, theta=0, sigma=4):
    np.random.seed(SEED)
    i = 0
    count = 0
    samples = np.empty(N)
    while i < N:
        x = np.random.normal(theta, sigma, 1)[0]
        u = np.random.uniform(0,1,1)[0]
        q_x = px(x, theta, sigma)
        p_x = pg(x, 1, theta, sigma)*q_x
        if u < p_x/(M*q_x):
            samples[i] = x
            i += 1
        count += 1
    acceptance = N / count
    return samples, acceptance

def importance_sampler(N, seed):
    return smc(0, N, seed)

def p_theta(theta, x=1.7, g=0, sigma=4):
    z = 1 + (theta/10)**2
    result = px(x, theta, sigma) * pg(x, g, theta, sigma)
    return result * 1 / (10 * np.pi * z)

def metro_hast(prop_sigma, N, seed, sigma=4, x=1.7, g=1):
    np.random.seed(seed)
    t = np.random.uniform(-20,20,1)[0]
    i = 0

```

```

count = 0
thetas = np.empty(N)
while i < N:
    tnew = np.random.normal(t, prop_sigma, 1)[0]
    gtnew_t = st.norm.pdf(tnew, loc=t, scale=prop_sigma)
    gt_tnew = st.norm.pdf(t, loc=tnew, scale=prop_sigma)
    pt = p_theta(t, x, g)
    ptnew = p_theta(tnew, x, g)
    accept = min(1, (ptnew * gt_tnew)/(pt * gtnew_t))
    u = np.random.uniform(0,1,1)[0]
    if u < accept:
        t = tnew
        thetas[i] = t
        i += 1
    count += 1
return thetas, acceptance

if __name__ == "__main__":
    theta = 0
    sigma = 4
    N = 10000
    SEED = 2019

# 1 a)
# Quadrature (Direct Integration)
x = np.linspace(start=-20, stop=20, num=N)
f = lambda x: px(x) * pg(x, 0)

p_absorb = scipy.integrate.quadrature(f, -20, 20)[0]
plt.plot(x, f(x))
plt.show()

print('Quadrature estimate of fraction of photons absorbed', \
      'average is %.3f' % (p_absorb))

# 1 b)
samples, acceptance = rejection_sampler(1, 10000, SEED)

sns.distplot(samples, kde = False)
plt.show()
print('Rejection Sampling: ')
print('Fraction of accepted sample is %.3f' % (acceptance))

# 1 c)

```

```

p_absorb_marginal = importance_sampler(1000, SEED)
print('Importance Sampling:')
print('Estimate of fraction of photons that get absorbed is %.3f' \
      % (p_absorb_marginal))

# 1 d)
x = 1.7
g = 1
thetas = np.linspace(start=-20, stop=20, num=N)
p_joint = p_theta(thetas, x, g)
plt.plot(thetas, p_joint)
plt.show()

# 1 e)
sample_t, acceptance = metro_hast(3, 50000, SEED)
sns.distplot(sample_t, kde = False)
plt.show()

print('Metropolis Hasting to sample theta:')
print('Fraction of accepted sample is %.3f' % (acceptance))

# 1 f)
p_abs_three = len(sample_t[(sample_t<3) & (sample_t>-3)]) \
              / len(sample_t)

print('Posterior probability of theta between' \
      ' +3 and -3 is %.3f' % (p_abs_three))

```